

Navigation Report

I used Deep Reinforcement Learning to solve this Banana Navigation Environment. I implemented Double DQN to generate the final trained agent.

Learning Algorithm

The final agent is trained with Double DQN (Deep Q-Network).

The environment has a state space of 37 dimensions and many of them are continuous. If we use a Q-learning algorithm, the memory and number of steps needed to explore all possibilities are not feasible. Instead, we can replace the Q-table with a neural network. Therefore, DQN is picked for this environment. Furthermore, to address the potential issue of overestimating action values, Double DQN is applied. The idea is to use two different networks to reduce the correlation between action selection and Q value evaluation. We use the DQN network to pick the next action, while use a target network to get the Q-value. In this way, only the “real” optimal option will be chosen and yield high reward. As a result, I implemented Double DQN to solve this environment.

Some hyperparameters used:

- Experience replay buffer size: 100000
- Training batch size: 64
- Reward discounting Gamma: 0.99
- Soft update Tau: 0.001
- Learning rate: 0.0005
- Update the weights every 4 steps
- Epsilon range: 1 to 0.01
- Epsilon decay factor: 0.99

Furthermore, the agent also supports DQN in its `learn()` function. Simply specify `double_dqn=False` in the `learn()` function to use DQN instead. In the early stage, I used DQN to solve the environment. But Double DQN significantly reduce the number of episodes needed to solve the environment and show more stable performance.

For the Q-Network, the architecture is a simple feed-forward network, with 3 Fully Connected layers and ReLu Activation Function. The details are:

```
QNetwork(  
    (fc1): Linear(in_features=37, out_features=64, bias=True)  
    (fc2): Linear(in_features=64, out_features=64, bias=True)  
    (fc3): Linear(in_features=64, out_features=4, bias=True)  
)
```

Model Performance

The trained agent can solve the environment within 400 episodes. The average 100-episode reward is about +16 after 1000 episodes. Here are the details:

```
Episode 100      Average Score: 2.33
Episode 200      Average Score: 6.66
Episode 300      Average Score: 10.99
Episode 350      Average Score: 13.05
  * Environment first solved in 350 episodes! Average
Score: 13.05. Continue training...
Episode 400      Average Score: 13.49
Episode 500      Average Score: 15.40
Episode 600      Average Score: 15.32
Episode 700      Average Score: 16.78
Episode 800      Average Score: 16.19
Episode 900      Average Score: 15.98
Episode 1000     Average Score: 16.18
```



Ideas for Future Work

From the related research in DQN, I think the following 2 techniques can help further improve the agent's performance:

- Prioritized experience replay:
 - There can be some cases that the TD error is high, and we should try to focus reducing such errors by giving it a higher weight in being selected from the experience replay
 - This should be able to help us use the past experiences more efficiently
- Dueling DQN
 - Separating state value estimates and advantage value estimates in the same network can potentially further help improve the model performance